

Finding the Needle in the Heap: Combining Static Analysis and Dynamic Symbolic Execution to Trigger Use-After-Free

Josselin Feist[†]
Laurent Mounier[†]
Sébastien Bardin[‡]
Robin David[‡]
Marie-Laure Potet[†]

December 2016

[†] Verimag (France), first.last@imag.fr

[‡] CEA LIST (France), first.last@cea.fr

Triggering *Use-After-Free*

Summary of this Presentation

- What is *Use-After-Free*
- Background on static analysis detecting *Use-After-Free* and dynamic symbolic execution (DSE)
- How to combine static analysis with DSE to trigger *Use-After-Free*
- Details for real-world application


Plan

- 1 Big Picture
- 2 Background
 - Use-after-free
 - Static Analysis Detecting UAF
 - Dynamic Symbolic Execution
- 3 Guided DSE
 - Intuition
 - Weighed-Slice
 - Oracle
 - Libraries
 - Implementation
- 4 Conclusion & Perspectives

Big Picture

Our Approach

- Use **static analysis** to extract vulnerable paths as slice of the program
- Create PoC confirming these results with **dynamic symbolic execution**



Techniques Combination

	Static analysis	DSE	Static analysis + DSE
Finding <i>Use-After-Free</i>	+	-	+
PoC generation	-	+	+

Advantages

- Allow unsound static analysis
- Complex *Use-After-Free* detection + PoC generation
- Correctness (if *Use-After-Free* found, true positive)

Plan

- 1 Big Picture
- 2 Background
 - Use-after-free
 - Static Analysis Detecting UAF
 - Dynamic Symbolic Execution
- 3 Guided DSE
 - Intuition
 - Weighed-Slice
 - Oracle
 - Libraries
 - Implementation
- 4 Conclusion & Perspectives

Plan

- 1 Big Picture
- 2 Background
 - Use-after-free
 - Static Analysis Detecting UAF
 - Dynamic Symbolic Execution
- 3 Guided DSE
 - Intuition
 - Weigthed-Slice
 - Oracle
 - Libraries
 - Implementation
- 4 Conclusion & Perspectives

Use-After-Free

Dangling Pointer

A dangling pointer is a pointer pointing to a free block, or to a block reallocated to another pointer.

```
1 | login=(char*) malloc(...);  
2 | ...  
3 | free(login); // login is now a dangling pointer
```


Use-After-Free

Dangling Pointer

A dangling pointer is a pointer pointing to a free block, or to a block reallocated to another pointer.

```
1 login=(char*)malloc(...);  
2 ...  
3 free(login); // login is now a dangling pointer
```

Use-After-Free

Use of a dangling pointer

```
1 login=(char*)malloc(...);  
2 ...  
3 free(login); // login is now a dangling pointer  
4 ..  
5 printf("%s\n", login) // use-after-free
```

Use-After-Free

Why it is dangerous?

- Re-allocations → create *indirect* aliases
- Information leakage, control-flow hijacking, ...

```
1 login=(char*)malloc(...);
2 ...
3 free(login); // login is now a dangling pointer
4 ..
5 passwords=(char*)malloc(...); // may re-allocate memory area used by login
6 ..
7 printf("%s\n", login) // prints the passwords !
```

Use-After-Free in the Wild

Particularities

- Difficult to detect (events distant, reasoning with heap, ..)
- No easy "pattern" (like for buffer overflow / string format)
- Lots of *Use-After-Free* in browsers
- Present in other apps (proftpd CVE-2011-4130, privoxy CVE-2015-1031, openssh...)

Detecting *Use-After-Free*

Dynamic Analysis

- Good dynamic detection (AddressSanitizer) / mitigation (MemGC in IE,...)
- Limitations:
 - Need to find the good path
 - Overhead during execution

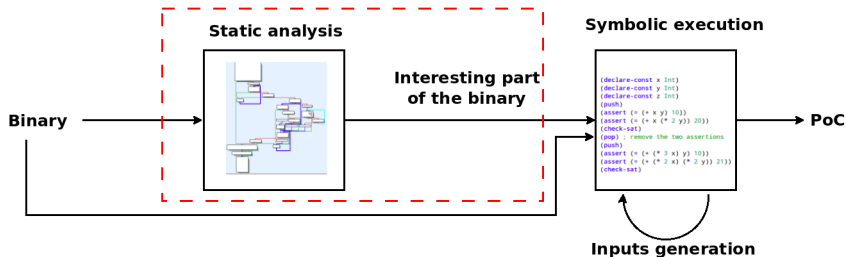
Static Analysis

- Source code: codesonar, coverity , ..
- Binary code: codesonar x86, veracode, ..
- Limitations:
 - Number of false positives
 - No PoC

Plan

- 1 Big Picture
- 2 Background
 - Use-after-free
 - Static Analysis Detecting UAF
 - Dynamic Symbolic Execution
- 3 Guided DSE
 - Intuition
 - Weigthed-Slice
 - Oracle
 - Libraries
 - Implementation
- 4 Conclusion & Perspectives


Static Analysis



Static analyzer: GUEB

Graph of Use-after-free Extraction from Binary

- Based on value analysis
- Open source: <https://github.com/montyly/gueb>
- Ocaml

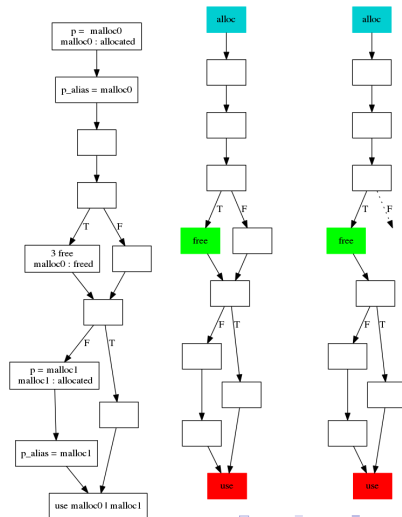


GUEB: Example

```

1 p=malloc(sizeof(int));
2 p_alias=p; // p and p_alias points
3           // to the same addr
4 read(f,buf,255); // buf is tainted
5
6 if(strncmp(buf,"BAD\n",4)==0){
7     free(p);
8     // exit() is missing
9 }
10 else{
11     .. // some computation
12 }
13
14 if(strncmp(&buf[4],"is a uaf\n",9)
15     ==0){
16     p=malloc(sizeof(int));
17 }
18 else{
19     p=malloc(sizeof(int));
20     p_alias=p;
21 }
22 *p=42; // not a uaf
23 *p_alias=43; // uaf if 6 and 14 =
                true

```



GUEB: Results

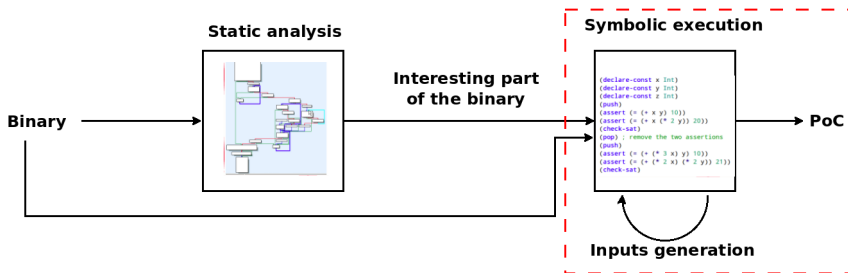
Results

- Several **new** *Use-After-Free* found (Jasper-JPEG-200 (CVE-2015-5221), openjpeg (CVE-2015-8871), gnome-nettool, accel-ppp, giflib (CVE-2016-3177), alsbat)
- Some results with good accuracy: gnome-nettool, 5 reports, just 4 false positives
- Some more difficult: jasper, 300+ reports
- Good results, but still need an expert to remove false positives
- Could we automatize the sorting of results?

Plan

- 1 Big Picture
- 2 Background
 - Use-after-free
 - Static Analysis Detecting UAF
 - Dynamic Symbolic Execution
- 3 Guided DSE
 - Intuition
 - Weigthed-Slice
 - Oracle
 - Libraries
 - Implementation
- 4 Conclusion & Perspectives

Dynamic Symbolic Execution



Dynamic Symbolic Execution

- Also called Whitebox fuzzing / Concolic Execution
- Program exploration technique
- Goal: generating inputs

DSE: Example

```
1 void f(int a){  
2   a = a+1;  
3   if(a == 0x42){  
4     printf("Win!\n");  
5   }  
6 }
```

Using Classical Fuzzing

- Hard to trigger randomly

DSE: Example

```
1 void f(int a){  
2   a = a+1;  
3   if(a == 0x42){  
4     printf("Win!\n");  
5   }  
6 }
```

Using Dynamic Symbolic Execution

- Choose seed (e.g.: $a=0$) → Generate trace 1-2-3-5-6
- Make `a` symbolic and generate path predicate

DSE: Example

```

1 void f(int a)
2   a = a+1;
3   if(a == 0x42){
4     printf("Win!\n");
5   }
6 }

```

Code	Concrete Values	Path Predicate
1 : void f(int a){	$a = 0$	a_0
2 : a = a+1;		
3 : if(a == 0x42){		

DSE: Example

```

1 void f(int a){
2   a = a+1;
3   if(a == 0x42){
4     printf("Win!\n");
5   }
6 }

```

Code	Concrete Values	Path Predicate
1 : <code>void f(int a){</code>	$a = 0$	a_0
2 : <code>a = a+1;</code>	$a = 1$	$a_1 = a_0 + 1$
3 : <code>if(a == 0x42){</code>		

DSE: Example

```

1 void f(int a){
2   a = a+1;
3   if(a == 0x42)
4     printf("Win!\n");
5 }
6 }

```

Code	Concrete Values	Path Predicate
1 : void f(int a){	$a = 0$	a_0
2 : a = a+1;	$a = 1$	$a_1 = a_0 + 1$
3 : if(a == 0x42){	$a = 1$	$a_1 = a_0 + 1 \wedge a_1 \neq 0x42$

DSE: Example

```

1 void f(int a){
2   a = a+1;
3   if(a == 0x42)
4     printf("Win!\n");
5 }
6 }

```

- Path predicate: invert conditional branch
→ explore new path
- Use SMT Solver (Z3, Boolector,..) to solve it
- Here, SAT, $a_0 = 0x41$
- Generate new trace with this new input

Code	Concrete Values	Path Predicate
1 : void f(int a){	$a = 0$	a_0
2 : a = a+1;	$a = 1$	$a_1 = a_0 + 1$
3 : if(a == 0x42){	$a = 1$	$a_1 = a_0 + 1 \wedge a_1 \neq 0x42$

DSE

- Large interest these past years
- Academic & Industrial interest
- SAGE, KLEE, S2E, Mayhem, Angr, Triton..
- Young topic, still a lot of limitations (and tuning needed)
- → Path explosion problem
- One key feature → How to prioritize paths exploration?

Instead of coverage → triggering specific path(s)

Plan

- 1 Big Picture
- 2 Background
 - Use-after-free
 - Static Analysis Detecting UAF
 - Dynamic Symbolic Execution
- 3 Guided DSE
 - Intuition
 - Weighed-Slice
 - Oracle
 - Libraries
 - Implementation
- 4 Conclusion & Perspectives

Plan

- 1 Big Picture
- 2 Background
 - Use-after-free
 - Static Analysis Detecting UAF
 - Dynamic Symbolic Execution
- 3 Guided DSE
 - Intuition
 - Weighed-Slice
 - Oracle
 - Libraries
 - Implementation
- 4 Conclusion & Perspectives

Guided DSE


Using Results from Static Analysis

- Explore only a sub-part of the program using GUEB

```

1  p=malloc(sizeof(int));
2  p_alias=p; // p and p_alias points
3             // to the same addr
4  read(f,buf,255); // buf is tainted
5
6  if(strncmp(buf,"BAD\n",4)==0){
7      free(p);
8      // exit() is missing
9  }
10 else{
11     .. // some computation
12 }
13
14 if(strncmp(&buf[4],"is a uaf\n",9)==0){
15     p=malloc(sizeof(int));
16 }
17 else{
18     p=malloc(sizeof(int));
19     p_alias=p;
20 }
21
22 *p=42 ; // not a uaf
23 *p_alias=43 ; // uaf if 6 and 14 = true

```



Guided DSE: Example


Example

- First input: filled with 'A'
- All condition of the trace except one go out of the subgraph
- Formula: $input[0..3] \neq "BAD \backslash n"$

```

1 | read(f,buf,255); // buf is tainted
2 |
3 | if(strncmp(buf,"BAD\n",4) == 0){ free(p);}
4 | else{ ..} // some computation

```



Guided DSE: Example


Example

- Second input: "BAD\n"
- No UaF
- Formula: $input[0..3] = \text{"BAD\\n"} \wedge input[4..13] \neq \text{"is a uaf\\n"}$

```

1 read(f,buf,255); // buf is tainted
2
3 if(strncmp(buf,"BAD\n",4) == 0){ free(p);}
4 if(strncmp(&buf[4],"is a uaf\n",9) == 0){
5 else{ ...}

```



Guided DSE: Example


Example

- Third input: "BAD\nis a uaf\n"
- UaF Found !

```

1 read(f,buf,255); // buf is tainted
2
3 if(strncmp(buf,"BAD\n",4) == 0){ free(p);}
4 if(strncmp(&buf[4],"is a uaf\n",9) == 0)
5
6 *p_alias = 42 ;

```



Plan

- 1 Big Picture
- 2 Background
 - Use-after-free
 - Static Analysis Detecting UAF
 - Dynamic Symbolic Execution
- 3 Guided DSE
 - Intuition
 - Weighed-Slice
 - Oracle
 - Libraries
 - Implementation
- 4 Conclusion & Perspectives

Weighed-Slice Guided DSE

WS-Guided

- Slice leads the exploration
- If a node is outside the slice → do not explore
- If a node is inside the slice → use score to prioritize
- Score: distance metric (i.e.: shortest path)
- Compute score as pre-preprocessing of the exploration
- Slice is *weighted* by the score

WS-Guided DSE


WS-Guided

- 3 events (alloc/free/use) → 3 different scores
- For a node, select score according past events on the trace

```

0 | p=malloc(..);
1 | for(..){
2 |   if(cond1)
3 |     free(p);
4 |   else
5 |     ..
6 | }
7 | *p=42;

```



Plan

- 1 Big Picture
- 2 Background
 - Use-after-free
 - Static Analysis Detecting UAF
 - Dynamic Symbolic Execution
- 3 Guided DSE
 - Intuition
 - Weighed-Slice
 - Oracle
 - Libraries
 - Implementation
- 4 Conclusion & Perspectives

Oracle

Use-After-Free in a trace?

- *Use-After-Free* detection on a trace is not so easy
- Trouble with aliases

```
1 int *p=malloc();
2 p_alias=p;
3 free(p);
4 if(cond){
5     p=malloc()
6 }
7 else{
8     p=malloc();
9     p_alias=p;
10 }
11 *p=0; // never uaf
12 *p_alias=0; // uaf if cond
```

Two paths (according *cond*), yet, *p* is always equal to *p_alias*.
Only *Use-After-Free* in *p_alias* if (*cond*)

Oracle

Use-After-Free in a trace?

- (i) $t = (\dots, n_{alloc}(size_{alloc}), \dots, n_{free}(a_f), \dots, n_{use}(a_u))$
- (ii) a_f is a reaching definition of the block returned by n_{alloc}
- (iii) a_u is a reaching definition of an address in the block returned by n_{alloc}

SMT-based solution

- Could use data dependencies analysis but:
 - Traces incomplet, need to have stubs for data dependencies
- Data dependencies kept *implicitly* with path predicate

Oracle

SMT-based solution

- New path predicate φ , *malloc* returns S_{alloc} , inputs conc
- $\Phi' = (a_f \neq S_{alloc}) \vee (a_u \notin [S_{alloc}, S_{alloc} + size_{alloc} - 1])$
- Oracle: $\varphi \wedge \Phi'$ is UNSAT \rightarrow *Use-After-Free*

Explanation

- Instead of $(\forall X \rightarrow SAT)$ we use $(\exists \overline{X} \rightarrow UNSAT)$
- $a_f \neq S_{alloc}$: the pointer given as the parameter for *free* is not the one allocated at n_{alloc} (negation of property (ii))
- $a_u \notin [S_{alloc}, S_{alloc} + size_{alloc} - 1]$: the pointer used is not a reaching definition of the pointer allocated at n_{alloc} (negation of property (iii))

Oracle

```

1 | int *p=malloc(4);
2 | p_alias=p;
3 | free(p);
4 | if(cond){
5 |   p=malloc()
6 | }
7 | else{
8 |   p=malloc();
9 |   p_alias=p;
10 | }
11 | *p=0; // never uaf
12 | *p_alias=0; // uaf if cond

```

First path

- Path predicate:

$$\varphi = (p_0 = S_{alloc} \wedge p_alias_0 = p_0 \wedge p_1 = 0x8040000)$$
- $\Phi' = (p_0 \neq S_{alloc} \vee p_alias_0 \notin [S_{alloc}, S_{alloc} + 3])$
- $\varphi \wedge \Phi'$ UNSAT: there is *Use-After-Free*

Oracle

```

1 | int *p=malloc(4);
2 | p_alias=p;
3 | free(p);
4 | if(cond){
5 |   p=malloc()
6 | }
7 | else{
8 |   p=malloc();
9 |   p_alias=p;
10 | }
11 | *p=0; // never uaf
12 | *p_alias=0; // uaf if cond

```

Second path

- Path predicate: $\varphi = (p_0 = S_{alloc} \wedge p_alias_0 = p_0 \wedge p_1 = 0x8040000) \wedge p_alias_1 = p_1)$
- $\Phi' = (p_1 \neq S_{alloc} \vee p_alias_1 \notin [S_{alloc}, S_{alloc} + 3])$
- $\varphi \wedge \Phi'$ SAT (e.g.: $S_{alloc} = 0$), no *Use-After-Free*

Plan

- 1 Big Picture
- 2 Background
 - Use-after-free
 - Static Analysis Detecting UAF
 - Dynamic Symbolic Execution
- 3 Guided DSE
 - Intuition
 - Weigthed-Slice
 - Oracle
 - Libraries
 - Implementation
- 4 Conclusion & Perspectives

Libraries

Handling libraries

- When apply to real-world, libraries are widely present
- Do not want to explore all of them
- We use two solutions:
 - Stubs → model effects on path predicate without tracing instructions (e.g.: if *realloc* returns new pointer, it performs copy of data)
 - Library Driven Heuristics (LDH) → known-behavior used to improve guiding

Library Driven Heuristics (LDH): Example

```
1 | p=malloc(size) ;  
2 | if(p == NULL)  
3 | {  
4 |     // path to trigger  
5 | }
```

Example

- To trigger the path → *malloc* needs to return 0

malloc heuristic

- Principle: on allocation functions, try as parameter a large number

Library Driven Heuristics (LDH): Example

```
1 read(f,tmp,255);
2 for(i=0;i<255;i++){
3     if(tmp[i]=='\0') break;
4     buf[i] = tmp[i];
5 }
6 buf[i]='\0';
7 if(strcmp(buf,"this is really bad") == 0)
8     ..
```

Example

- Comparison on string whose length depends of loop iteration number
- Here, every time loop iterates: $buf[i] \neq '\0'$
- To solve *strcmp* condition, need a trace unrolling 18 times the loop

Library Driven Heuristics (LDH): Example

```
1 read(f,tmp,255);
2 for(i=0;i<255;i++){
3     if(tmp[i]=='\0') break;
4     buf[i] = tmp[i];
5 }
6 buf[i]='\0';
7 if(strcmp(buf,"this is really bad") == 0)
8     ..
```

strcmp heuristic

- Principle: on this pattern, use size of constant strings passed to `strcmp` to find the desired iteration.

Plan

- 1 Big Picture
- 2 Background
 - Use-after-free
 - Static Analysis Detecting UAF
 - Dynamic Symbolic Execution
- 3 Guided DSE
 - Intuition
 - Weigthed-Slice
 - Oracle
 - Libraries
 - Implementation
- 4 Conclusion & Perspectives

Guided DSE

BINSEC/SE

- Open source platform to perform static analysis and DSE ^a
 - Guiding module based on Ocaml functors
 - Concretization / Symbolization Policies
 - Other nice features (opaque predicates, ..)
- Ocaml
- <http://binsec.gforge.inria.fr/tools>

^aCode of this talk available in the next release (this month)

Guided DSE

Demo Binsec (Jasper-JPEG-200 CVE-2015-5221)

Jasper

Details on DSE Exploration

- 20 mins
- 9 test cases generated
- Last test case triggering the *Use-After-Free*
- More details in the paper (comparison with fuzzing)

Plan

- 1 Big Picture
- 2 Background
 - Use-after-free
 - Static Analysis Detecting UAF
 - Dynamic Symbolic Execution
- 3 Guided DSE
 - Intuition
 - Weighed-Slice
 - Oracle
 - Libraries
 - Implementation
- 4 Conclusion & Perspectives

Conclusion

Guided DSE

- First *end-to-end* approach targeting *Use-After-Free*
- Operates on binaries
- Working on real CVE
- Everything is open source
- Combining program analysis techniques → way to go

Limitation

- More complex software (browser, multi-thread)..
• Scalability issue ?

Perspectives

Guided DSE

- Improve guiding
- Other score (data-flow?)
- Need more benchmarks (network, etc...)
- Combine with fuzzing
- Other applications (BinDiff, 1day generation, ..)

XP

```
MIF
component
```

Figure: PoC of CVE-2015-5221 (test.plain)

```
jasper --input test.plain --input-format mif --output out --output-format mif
```

XP

Name	Time	MIF line	UAF found	# Paths
DSE (in BINSEC/SE)				
WS-Guided+LDH	20m	3min	Yes	9
WS-Guided	6h	3min	No	44
DFS(slice)	6h	3min	No	68
DFS	6h	3min	No	354
standard fuzzers (arbitrary seed)				
AFL	7h	< 1min	No	174 [†]
Radamsa	7h	> 1h	No	~ 1000000 [‡]
standard fuzzers (MIF seed)				
AFL (MIF input)	< 1min	< 1min	Yes	< 10
Radamsa (MIF input)	< 1min	< 1min	Yes	< 10

[†] AFL generates more input, 174 is the number of unique paths.

[‡] For radamsa it is not trivial to count the number of unique paths.