

Finding the Needle in the Heap: Combining Binary Analysis Techniques to Trigger Use-After-Free

Josselin Feist

Supervisors: Marie-Laure Potet and Laurent Mounier

March 29, 2017



Plan

- 1 Context: Vulnerabilities
 - Program Vulnerabilities
 - Automated Program Analysis
- 2 Use-After-Free, a Complex Vulnerability
- 3 Using Static Analysis to Detect Use-After-Free
- 4 Using Dynamic Symbolic Execution to Trigger Use-After-Free
- 5 Thesis Contributions and Conclusion

Context

Security

- Recent interest from media and general public for computer security;
- Topic: security of computer programs
 - Programs contain bugs;
 - Bugs range from benign (almost no impact) to critical;
 - Some bugs decrease the security of the system → called **vulnerabilities**.
- Dirty Cow [Goo], Stagefright [Fin], Heartbleed [McM], ...

Vulnerabilities

Different categories of vulnerabilities

- Bad use of cryptography;
- Unsanitized inputs, such as SQL Injection;
- **Memory corruptions** ~ low-level vulnerabilities.

Possible consequences of memory corruption

- The system becomes unavailable (e.g., *denial-of-service* attacks);
- Critical information is leaked (such as cryptographic keys);
- The full system is compromised (unauthorized code execution).

Who Has Interest in Vulnerabilities?

Who uses them?

- Malicious individuals / groups (malware market, ...)
- Governments, industrial theft, ...

Who finds them?

- Developers or internal researchers of vendors;
- *Bug-bounties* programs;
- Black market → prices can reach millions of dollars.

Practical topic

Vulnerabilities research requires engineering process:

- Need for **working solutions**;
- **PoC** and **reproducibility** are important.

Finding Vulnerabilities

Manual or **assisted** by automated program analysis techniques

Dynamic analysis

- Fuzzing: generating inputs stressing the program (AFL, Radamsa, ..);
- Simple but efficient → **widely used** in industry;
- Not well-adapted to detect complex patterns.

Static analysis

- Analyzing several behaviors of a program without executing it;
- Historically used for verification (proof the absence of bugs);
- **Less used** in security community (number of **false alarms**, no PoC).

Other techniques

Distinction between dynamic and static not always pertinent:

- Symbolic execution (SAGE [GLM12], Mayhem [CARB12], ..);
- Guided fuzzing (Libfuzzer);
- Combining techniques [ZC10, BMMS11, HSNB13, BCDK14].

Need for binary analysis

- Source code not available;
- Undefined behaviors;
- Precise memory layout.

→ Binaries analysis is harder than source code analysis, but it is mandatory in several contexts.

Vulnerability Research: Past, Present, Future

Past

- Buffer overflow, string format;
- Lots of manual efforts to find and exploit vulnerabilities;
- Individual hackers, few professionals.

Present

- Overflow hard to exploit → use-after-free, type confusion;
- Better tools: *Smart* fuzzers, etc.;
- Business model of entire companies.

Future

- Darpa CGC → fully-automated system;
- New attack vectors (Rowhammer [SD15]);
- Looking for protection killing class of vuln. or exploits.

Plan

- 1 Context: Vulnerabilities
- 2 Use-After-Free, a Complex Vulnerability
 - Use-After-Free Description
 - Use-After-Free Detection
- 3 Using Static Analysis to Detect Use-After-Free
- 4 Using Dynamic Symbolic Execution to Trigger Use-After-Free
- 5 Thesis Contributions and Conclusion

Use-After-Free: Example

```
1 char *login, *password;  
2 login = malloc(..);  
3 ...  
4 free(login);  
5 // login still points to the address returned by malloc
```




Figure: At line 2





Figure: At line 5


Use-After-Free: Example

```
1 char *login, *password;  
2 login = malloc(..);  
3 ...  
4 free(login);  
5 // login still points to the address returned by malloc  
6 password = malloc(..);
```



Use-After-Free: Example

```
1 char *login, *password;  
2 login = malloc(..);  
3 ...  
4 free(login);  
5 // login still points to the address returned by malloc  
6 password = malloc(..);  
7 printf("Login: %s\n", login); // prints the password
```



Definitions

Dangling pointer

Pointer referencing a free block or a block reallocated to another pointer.

```
1 char *login;  
2 login = malloc(...);  
3 free(login);
```

Use-After-Free

Use of a dangling pointer.

```
1 char *login;  
2 login = malloc(...);  
3 free(login);  
4 [...];  
5 printf("..", login);
```

Dangerousness

- Reallocation of memory area used by p ?
- Can lead to code execution, ...

A Recent Vulnerability

- Recent vulnerability → less studied;
- Demonstrated to be dangerous: Operation Aurora [Wik13], Pwn2Own, ...




Figure: CVE: number of Use-After-Free (2016-12-28).

Particularities of Use-After-Free

- Three events:
 - Allocation,
 - Free,
 - Use;
- Events can be distant in the code → need for **scalability**;
- Pointers are complex to track (e.g., *aliases* problem, ...) → need for **precision**;
- No easy "pattern" (like for buffer overflow / string format).

Detecting Use-After-Free – State of the Art Methods

Problems with existing methods

- Not sufficient alone
 - Dynamic detectors, such as ASan [SBPV12], Valgrind [Val], need other techniques to trigger the path;
- Not applicable in a security perspective
 - Static analyzers with too many false alarms and no PoC;
 - Need for manual annotations;
- Academic papers without tools available → no fair comparison;
- Industrial tools without description.

Thesis Goals and Challenges

- **Apply formal methods to security purpose;**
- Develop techniques precise enough to **detect** and **trigger Use-After-Free**;
- Apply static analysis to **real-world binaries**:
 - Scalable analysis;
 - Programs designed without security in mind;
 - Avoiding no realistic working hypotheses.

→ Finding which methods can be applied in real contexts;
→ Finding the right trade-off between precision and scalability.

Thesis Results

- Developing a static analysis unsound, but well-suited to detect **Use-After-Free on binaries**;
- Using dynamic symbolic execution to **remove false alarms** of the static analysis and to **generate PoC**;
- Found several **previously unknown** Use-After-Free in software;
- Open-source tool-chain.





Figure: Architecture of our approach.

Plan

- 1 Context: Vulnerabilities
- 2 Use-After-Free, a Complex Vulnerability
- 3 Using Static Analysis to Detect Use-After-Free
 - Value Set Analysis
 - Detecting Use-After-Free
 - Real-World Binary Code Analysis
 - Static Analysis: Conclusion
- 4 Using Dynamic Symbolic Execution to Trigger Use-After-Free
- 5 Thesis Contributions and Conclusion



Statically Detecting Use-After-Free

Steps

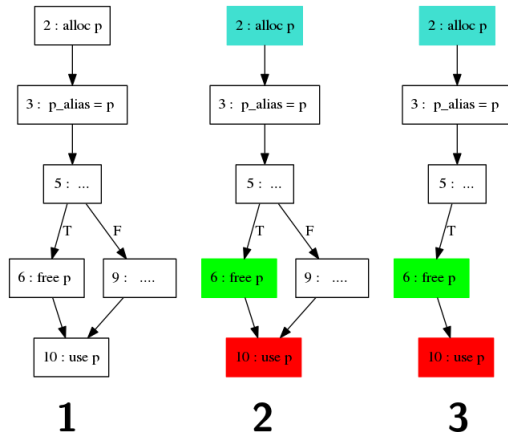
- 1 Value analysis → tracking use of pointers,
- 2 Characterization of Use-After-Free,
- 3 Slice extraction.

The Three Steps

```

1  ...
2  p=malloc(..);
3  p_alias=p;
4  ..
5  if(..){
6    free(p);
7  }
8  else{ ..}
9  ..
10 *p_alias = 42;

```



Value Set Analysis (VSA)

Binary code = no variables, only memory and registers accesses.

VSA: background [BR10]

- For each point of program, represent all possible memory states;
- *Transfer functions* → the transitions between instructions;
- Loop / recursion → compute fixed points (*costly*).

```

1  a = 1;          1  mov eax,1
2  if(cond){      2  jne L1
3    a = 2;        3  mov eax,2
4  }              4  L1:
5  a = a + 1;      5  add eax,1

```

Inst	Memory state
1: mov eax,1	$\text{eax} \in [1]$
3: mov eax,2	$\text{eax} \in [2]$
5: add eax,1	$\text{eax} \in [2,3]$

Our Value Analysis

Detection requirements

- Track the use of pointers / aliases;
- Find the state of the heap objects (allocated / freed);
- Find paths leading to Use-After-Free.

Light VSA

- *Best-effort* to track values;
- Unroll loops and inline functions;
- One allocation = one **new** memory area.

Our Memory Model

Overview

- For each instruction \rightarrow **AbsEnv**;
- *AbsEnv* associates to each memory location **memLoc** a set of possible values **valueSet**;

- Example of *memLoc*:
 - *eax*,
 - $[esp_0 - 4]$: local variable,
 - $[chunk_0 + 0]$: heap chunk.
- Example of *valueSet*:
 - $chunk_0 + \{0, 4, 8\}$.




Figure: Local variables.

Memory Model: Example

```

1  p=0; // p is accessed through the initial value of esp0 -4
2  if(cond){
3    p=malloc(); // malloc returns chunk0 + {0}
4  }
5  *p = .. // p = malloc or p = 0

```




Figure: *AbsEnv* at line 5.

Memory Model: Allocation Status

How to represent the status (allocated / freed) of a chunk?

Three solutions proposed

- Object-Based: status kept apart of the memory model (using two functions: *HA* and *HF*);
- Pointer-Based: status kept into the *valueSet*;
- A last variant allows to detect *Stack-based* Use-After-Free.

Details

- Conservative free;
- Object-Based is more classic;
- Pointer-Based has better precision when path conditions are ignored.

Allocation Status: Limitation of Object-Based

Limitations when dealing with path conditions

```

1  int *p=malloc(sizeof(int));
2  if(cond){
3      free(p);
4      p=malloc(sizeof(int));
5  }
6  // union of memory states

```

Code	Init _{reg}	Heap State
1: <code>p=malloc()</code>	$(esp_0 - 0x4) \rightarrow \text{chunk}_0$	$HA = \{\text{chunk}_0\}$ $HF = \emptyset$
3: <code>free(p)</code>	$(esp_0 - 0x4) \rightarrow \text{chunk}_0$	$HA = \emptyset$ $HF = \{\text{chunk}_0\}$
4: <code>p=malloc()</code>	$(esp_0 - 0x4) \rightarrow \text{chunk}_1$	$HA = \{\text{chunk}_1\}$ $HF = \{\text{chunk}_0\}$
6: <code>// union</code>	$(esp_0 - 0x4) \rightarrow \text{chunk}_0, \text{chunk}_1$	$HA = \{\text{chunk}_1\}$ $HF = \{\text{chunk}_0\}$

Table: VSA results using the *object-based* representation.

Pointer-Based Representation

```

1  int *p=malloc(sizeof(int));
2  if(cond){
3    free(p);
4    p=malloc(sizeof(int));
5  }
6  // union of memory states

```

Code	Init _{reg}
1: <code>p=malloc()</code>	$(esp_0 - 0x4) \rightarrow (\text{chunk}_0, A)$
3: <code>free(p)</code>	$(esp_0 - 0x4) \rightarrow (\text{chunk}_0, F)$
4: <code>p=malloc()</code>	$(esp_0 - 0x4) \rightarrow (\text{chunk}_1, A)$
6: <code>// union</code>	$(esp_0 - 0x4) \rightarrow (\text{chunk}_0, A), (\text{chunk}_1, A)$

Table: VSA results using the *pointer-based* representation.

→ No dangling pointer is kept.

Allocation Status: Comparison

Discussion

- Benchmarks show that:
 - Pointer-based **reduce by two** the number of false alarms;
 - Pointer-based comes with no time or space overhead;
- If the analysis does not handle path conditions, pointer-based is to be used.

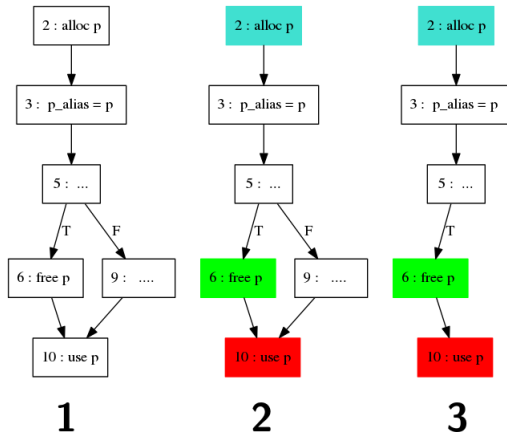
Use-After-Free Characterization and Detection

VSA results

- For each instruction \rightarrow the memory state;
- For each heap element \rightarrow allocated or freed;
- Use-After-Free characterization – checking the status of the heap elements on memory dereferencing:
 - load src, dst $\rightarrow mem(src)$ is freed ?
 - store src, dst $\rightarrow mem(dst)$ is freed ?
- A slice representing all paths going through the three events is extracted.

Detecting Use-After-Free

```
1  ...
2  p=malloc(..);
3  p_alias=p;
4  ..
5  if(..){
6    free(p);
7  }
8  else{ ..}
9  ..
10 *p_alias = 42;
```



Application to Real-World Binary Code

VSA handling real binary code

- Handling nested and irreducible loops;
- Dealing with errors in the CFG;
- Heuristics adapting the VSA:
 - Function without return statement,
 - Lost of the stack frame,
 - ...
- Discussion on the validity of our VSA:
 - Results **inconsistent** with reality, yet **well-suited in practice**.

Static Binary Analysis: Validity

Origin	Type
No handling of conditions	Over-approximation
Weak updates	Over-approximation
Unrolling (missing paths)	Under-approximation
Inlining bounded by size	Under-approximation
Inlining bounded by depth	Under-approximation
Aliases between uninitialized values	Under-approximation
Inlining bounded by depth	Inconsistency
Recursion	Inconsistency
Incorrect CFG	Inconsistency
No overlap between memory location	Inconsistency
Unrolling (invalid paths)	Inconsistency
Ignored updates	Inconsistency

Table: Summary of types of approximations for static analysis.

Suitable Results

Help analyzing the results

- A same dangling pointer can be used at multiples locations;
 - Solution: **grouping similar** Use-After-Free;
- Due to inlining, a function f containing a Use-After-Free is reported as many times it is called;
 - Solution: **signature to detect likely-similar** Use-After-Free.

Experiments

Implementation

- GUEB: <https://github.com/montyly/gueb>
- Ocaml
- Use with IDA/BinNavi (REIL as intermediate representation).

Questions

- Is GUEB precise enough to detect unknown Use-After-Free without raising too many false alarms?
- Is GUEB robust enough to be applied at scale?
- Two experimental results:
 - Finding previously unknown vulnerabilities in 6 software;
 - GUEB applied to a significant number of binaries (488).

Previously Unknown Use-After-Free Found

Name	#Lines	Time	#UAF	#Signature	#REIL ins.
alsabat	~ 2000	7s	1	1	99933
gnome-nettool	~ 6500	17s	7	5	260882
gifcolor (CVE-2016-3177)	~ 9000	21s	15	12	233303
jasper (CVE-2015-5221)	~ 34200	4m23	255	114	2154927
accel-ppd	~ 61000	5m5	35	30	3907862
openjpeg (CVE-2015-887)	~ 205200	6m10	329	300	2170081

Experiments

Discussion

- GUEB yields an *acceptable* number of results;
- Largest (300) takes time to analyze, but it still realistic;
- In practice iterative analyze, several false alarms are removed by adding user-provided stubs to GUEB.

Experiments: Robustness

Results

- 488 binaries from `/usr/bin` of Ubuntu 16.04,
- 406 without Use-After-Free found,
- 82 with Use-After-Free found. Most likely to be false alarms,
- Recall: too many false alarms is worse than missing Use-After-Free.

# Bin	Signature (\bar{x} , med, max)	UAF (\bar{x} , med, max)	Time Total (\bar{x} , med, max)
82	(9, 3, 210)	(16, 5, 247)	0h 52m 48s (38s, 27s, 4m 27s)

Static Analysis: Contributions

Contributions

- Memory model and VSA well adapted to binary code and scalability;
- Study of heap objects modeling and UaF detection;
- Implementation and benchmarks demonstrating its efficiency and robustness.

Limitations


- Too large programs,
- Multi-threading,
- Not well suited for some programs (e.g., *reference counters*).

→ An analyst is still needed. Could we automatize more?

Generating PoC?

Plan

- 1 Context: Vulnerabilities
- 2 Use-After-Free, a Complex Vulnerability
- 3 Using Static Analysis to Detect Use-After-Free
- 4 Using Dynamic Symbolic Execution to Trigger Use-After-Free
 - Dynamic Symbolic Execution: Background
 - Weighted-Slice Guided Dynamic Symbolic Execution
 - Implementation and Benchmark
 - Dynamic Symbolic Execution: Conclusion
- 5 Thesis Contributions and Conclusion



Dynamic Symbolic Execution (DSE)

DSE: an automated input generation technique.

How it works

- From a program and an input, generate an execution trace;
- Build a *path predicate* from the trace = logical formula representing the trace as a set of constraints over the inputs;
- Use SMT solver to invert conditional instruction;
- If SAT, generate new inputs.

DSE: Example

- First input: $a = 0$, generating a first trace;
- Build the path predicate;
- Invert condition at line 3 $\rightarrow a_0 \wedge a_1 = a_0 + 1 \wedge a_1 \neq 0x42$;
- Solve the formula: $a_0 = 0x41$, new input reaching line 4.

```

1 void f(int a){
2   a = a+1;
3   if(a == 0x42){
4     printf("Win!\n");
5   }
6 }

```

Trace

Path Predicate

1: int a

 $a_0 \wedge$

2: a = a + 1

 $a_1 = a_0 + 1 \wedge$

3: a == 0x42

 $a_1 \neq 0x42$

5: ..

DSE

Large recent interest in security

- Academic & Industrial interest;
- SAGE, KLEE, Mayhem, Angr, Triton, etc.;
- Young topic, still a lot of limitations;
- Other use: deobfuscation [DB16], etc.

Challenges for path exploration

- Path explosion problem:
 - → Importance of the **exploration strategy**;
- Tuning on path predicate:
 - Concretization,
 - Inputs functions,
 - Libraries.

Guided DSE

Guided DSE: Using a score function to prioritize the exploration.

Our approach: Weigthed-Slice

- Guided DSE toward the Use-After-Free slice;
- 3 events (alloc / free / use) \rightarrow 3 scores;
- DSE guided by the *next event to reach*;
- Partition nodes in a trace according if an event occurred.

We called our solution: *Weighted-Slice*.

Weighted-Slice Guided DSE: Example




Figure: Slice




Figure: Trace t

Weighted-Slice Guided DSE




Figure: Trace with partition




Figure: Trace with scores

Detecting Use-After-Free Without False Alarms

Detecting Use-After-Free on a trace

- Not so trivial;
- Trouble with indirect aliases.

```
1 p=malloc(..); // return X
2 free(p);      // free(X)
3 p2=malloc();  // return X
4 *p2 = 0;      // [X] = 0
5 *p = 0;       // [X] = 0
```

Solution

- Need data dependencies, allocator modifications or metadata;
- Proposition: Oracle based on symbolic execution.

Validity of the DSE

- DSE correctness: yes \rightarrow paths generated are feasible;
 - DSE completeness: no \rightarrow bounded exploration;
 - Oracle correctness: yes \rightarrow Use-After-Free detected are true positives;
 - Oracle completeness: no \rightarrow Use-After-Free can be not detected.
-
- We produce true positive without false positives;
 - Use-After-Free can be missed.

Experiment

Implementation

- Guided DSE implemented into the BINSEC/SE platform;
- Design of a generic exploration mechanism;
- Shortest path to the destination.

Questions

- Is DSE exploration working on real-world examples?
- Is the engine able to trigger Use-After-Free?
- Is GUEB helping the exploration to trigger the Use-After-Free?
- Experimental result:
 - The method triggers the Use-After-Free of the Jasper CVE.

Experiment

- Tested on JasPer, on a slice we knew to contain a Use-After-Free, with as seed a file filled with 'A';
- Comparison with standard exploration algorithm and fuzzers.

Name	Time	UAF found	# Paths
WS-Guided (gueb)	20m	Yes	9
<i>DFS (slice)</i>	6h	<i>No</i>	68
<i>DFS</i>	6h	<i>No</i>	354
<i>AFL</i>	7h	<i>No</i>	174
<i>Radamsa</i>	7h	<i>No</i>	<i>N/A</i>
<i>AFL (better seed file)</i>	< 1min	Yes	< 10
<i>Radamsa (better seed file)</i>	< 1min	Yes	< 10

Experiment: Discussion and Limitations

- Only our solution found the Use-After-Free without a seed;
- Fuzzers found the vulnerability with a proper seed (only few mutations are needed in this case);
- Promising preliminary results, but:
 - A larger experiment is needed to validate the approach;
 - BINSEC/SE is still young, only JasPer was running properly;
 - The slice containing the Use-After-Free was explored → all the slices (~ 200) have to be explored, but it is still realistic.

Dynamic Symbolic Execution: Conclusion

Contributions

- How to guide DSE toward Use-After-Free using GUEB;
- Creation of a *proof-of-concept* on JasPer;
- Implementation into the BINSEC platform;
- Tuning on DSE:
 - Initial memory state;
 - Exploration enhancement based on programming patterns.

Plan

- ① Context: Vulnerabilities
- ② Use-After-Free, a Complex Vulnerability
- ③ Using Static Analysis to Detect Use-After-Free
- ④ Using Dynamic Symbolic Execution to Trigger Use-After-Free
- ⑤ Thesis Contributions and Conclusion
 - Approach Discussion
 - Contributions and Perspectives




Figure: Architecture of our approach.

Combining Static Analysis with DSE: Validity of the results

GUEB	Results	BINSEC/SE	Results
Paths analyzed	Feasible, infeasible and unexplored paths	Concrete traces	All paths are feasible
		Bounded exploration	Unexplored paths
UaF detected	True positives and False positives	Oracle	True positives and no false positives
		Bounded exploration	False negatives
UaF not detected	False negatives	Not explored	False negatives

Table: Validity of the global approach.

Contributions

Static analysis

- Design of a memory model and a VSA well suited to be applied on real-world binary code;
- Study of the heap model;
- Use-After-Free characterization and representation to provide suitable results.

Dynamic symbolic execution

- Exploration algorithm using information provided by static analysis;
- Refinements on the DSE exploration and path predicates computation.

Contributions

Implementation and experiments

- Discovery of 6 new vulnerabilities;
- Creation of a PoC;
- All tools are open-source;
- All files are available to reproduce the experiments.

→ First **end-to-end** approach targeting Use-After-Free.

Perspectives

Static analysis

- Dedicated static analysis, less scalable, but more precise → reduce false alarms;
- Target other vulnerabilities (Use-Before-Initialization [LWP⁺17]), or specific software (e.g., kernel).

Dynamic symbolic execution

- Testing the guided DSE on other programs;
- Combining with fuzzers [SGS⁺16];
- Slices from other origins (e.g., BinDiff on patched programs).

Use-After-Free

- Study of the exploitability.

Thesis Publications

- Josselin Feist, Laurent Mounier, and Marie-Laure Potet. [Statically detecting Use-after-Free on binary code.](#)
2013 JCVHT;
- Josselin Feist, Laurent Mounier, and Marie-Laure Potet. [Using static analysis to detect use-after-free on binary code.](#)
In *1st Symposium on Digital Trust in Auvergne*, 2014;
- Josselin Feist. [Gueb : Static detection of use-after-free on binary.](#)
In *ToorCon San Diego*, 2015;
- Josselin Feist, Laurent Mounier, and Marie-Laure Potet. [Guided dynamic symbolic execution using subgraph control-flow information.](#)
In *SEFM*, Lecture Notes in Computer Science, pages 76–81. Springer, 2016;
- Josselin Feist, Laurent Mounier, Sébastien Bardin, Robin David, and Marie-Laure Potet. [Finding the needle in the heap: Combining static analysis and dynamic symbolic execution to trigger use-after-free.](#)
In *6th Software Security, Protection, and Reverse Engineering Workshop, SSPREW 2016, Los Angeles, CA, USA, December 5-6, 2016*, pages 1–12, 2016.

References I



Sebastien Bardin, Omar Chebaro, Mickal Delahaye, and Nikolai Kosmatov.

An all-in-one toolkit for automated white-box testing.

In Martina Seidl and Nikolai Tillmann, editors, *TAP*, volume 8570 of *Lecture Notes in Computer Science*, pages 53–60. Springer, 2014.



Domagoj Babic, Lorenzo Martignoni, Stephen McCamant, and Dawn Song.

Statically-directed dynamic automated test generation.

In Matthew B. Dwyer and Frank Tip, editors, *ISSTA*, pages 12–22. ACM, 2011.



Gogul Balakrishnan and Thomas Reps.

Wysinwyx: What you see is not what you execute.

ACM Trans. Program. Lang. Syst., 32(6):23:1–23:84, August 2010.



Sang Kil Cha, Thanassis Avgerinos, Alexandre Rebert, and David Brumley.

Unleashing mayhem on binary code.

In *Proceedings of the 2012 IEEE Symposium on Security and Privacy*, SP '12, pages 380–394, Washington, DC, USA, 2012. IEEE Computer Society.



Robin David and Sebastien Bardin.

Code deobfuscation: Intertwining dynamic, static and symbolic approaches.

Black Hat Europe, 2016.



Josselin Feist.

Gueb : Static detection of use-after-free on binary.

In *ToorCon San Diego*, 2015.

References II



Jon Fingas.

Stagefright exploit reliably attacks android phones.

<https://www.engadget.com/2016/03/19/reliable-stagefright-android-exploit/>.



Josselin Feist, Laurent Mounier, Sébastien Bardin, Robin David, and Marie-Laure Potet.

Finding the needle in the heap: Combining static analysis and dynamic symbolic execution to trigger use-after-free.

In *6th Software Security, Protection, and Reverse Engineering Workshop, SSPREW 2016, Los Angeles, CA, USA, December 5-6, 2016*, pages 1–12, 2016.



Josselin Feist, Laurent Mounier, and Marie-Laure Potet.

Using static analysis to detect use-after-free on binary code.

In *1st Symposium on Digital Trust in Auvergne*, 2014.



Josselin Feist, Laurent Mounier, and Marie-Laure Potet.

Guided dynamic symbolic execution using subgraph control-flow information.

In *SEFM, Lecture Notes in Computer Science*, pages 76–81. Springer, 2016.



Josselin Feist, Laurent Mounier, and Marie-Laure Potet.

Statically detecting Use-after-Free on binary code.

2013 JCVHT.



Patrice Godefroid, Michael Y. Levin, and David A. Molnar.

Sage: Whitebox fuzzing for security testing.

ACM Queue, 10(1):20, 2012.

References III



Dan Goodin.

most serious linux privilege-escalation bug ever is under active exploit.

<http://arstechnica.com/security/2016/10/>

[most-serious-linux-privilege-escalation-bug-ever-is-under-active-exploit/](http://arstechnica.com/security/2016/10/most-serious-linux-privilege-escalation-bug-ever-is-under-active-exploit/).



Istvan Haller, Asia Slowinska, Matthias Neugschwandtner, and Herbert Bos.

Dowsing for overflows: A guided fuzzer to find buffer boundary violations.

In *Proceedings of the 22Nd USENIX Conference on Security, SEC'13*, pages 49–64, Berkeley, CA, USA, 2013. USENIX Association.



Kangjie Lu, Marie-Therese Walter, David Pfaff, N Stefan, Wenke Lee, and Michael Backes.

Unleashing use-before-initialization vulnerabilities in the linux kernel using targeted stack spraying. NDSS, 2017.



Robert McMillan.

How heartbleed broke the internet and why it can happen again.

<https://www.wired.com/2014/04/heartbleedslesson/>.



Konstantin Serebryany, Derek Bruening, Alexander Potapenko, and Dmitry Vyukov.

Addresssanitizer: A fast address sanity checker.

In *USENIX ATC 2012*, 2012.



Mark Seaborn and Thomas Dullien.

Exploiting the dram rowhammer bug to gain kernel privileges.

Black Hat, 2015.

References IV



Nick Stephens, John Grosen, Christopher Salls, Andrew Dutcher, Ruoyu Wang, Jacopo Corbetta, Yan Shoshitaishvili, Christopher Kruegel, and Giovanni Vigna.

Driller: Augmenting fuzzing through selective symbolic execution.

In *NDSS*. The Internet Society, 2016.



Valgrind.

Memcheck.

<http://valgrind.org/info/tools.html>.



Wikipedia.

Operation aurora, 2013.

[Online; accessed 01-March-2013].



Cristian Zamfir and George Candea.

Execution synthesis: a technique for automated software debugging.

In Christine Morin and Gilles Muller, editors, *EuroSys*, pages 321–334. ACM, 2010.

memLoc

- The **constructor** determines the region of a *memLoc*;
- 6 regions:
 - **Globals** of *addr*
 - **Registers** of *reg_name*
 - **He** of *chunk* \times *offset* (holding heap elements)
 - **Init_{reg}** of *id* \times *offset* (holding initial values of registers)
 - **Init_{mem}** of *id* \times *offset* (holding initial values of the memory)
 - \top_{loc} .

valueSet

- The **constructor** determines the base of a *valueSet*;
- A *valueSet* = $\{base \times \{\mathbb{N}\}\}$, 4 bases:
 - **Constant**
 - **He** of *chunk*
 - **Init_{reg}** of *id*
 - **Init_{mem}** of *id*.

Oracle

Use-After-Free in a trace

- Use-After-Free detection on a trace is not so easy;
- Trouble with aliases.

```
1  int *p=malloc();
2  p_alias=p;
3  free(p);
4  if(cond){
5      p=malloc()
6  }
7  else{
8      p=malloc();
9      p_alias=p;
10 }
11 *p=0; // never uaf
12 *p_alias=0; // uaf if cond
```

Two paths (according *cond*), yet *p* is always equal to *p_alias*.
Only Use-After-Free in *p_alias* if (*cond*).

Oracle

Use-After-Free in a trace

- (i) $t = (\dots, n_{alloc}(size_{alloc}), \dots, n_{free}(a_f), \dots, n_{use}(a_u))$;
- (ii) a_f is a reaching definition of the block returned by n_{alloc} ;
- (iii) a_u is a reaching definition of an address in the block returned by n_{alloc} .

SMT-based solution

- Could use data dependencies analysis but:
 - Traces incomplete, need to have stubs for data dependencies;
- Data dependencies kept *implicitly* with path predicate.

Oracle

SMT-based solution

- New path predicate φ , *malloc* returns S_{alloc} , inputs *conc*;
- $\Phi' = (a_f \neq S_{alloc}) \vee (a_u \notin [S_{alloc}, S_{alloc} + size_{alloc} - 1])$;
- Oracle: $\varphi \wedge \Phi'$ is UNSAT \rightarrow Use-After-Free.

Explanation

- Instead of $(\forall X \rightarrow SAT)$ we use $(\exists \bar{X} \rightarrow UNSAT)$;
- $a_f \neq S_{alloc}$: the pointer given as the parameter for *free* is not the one allocated at n_{alloc} (negation of property (ii));
- $a_u \notin [S_{alloc}, S_{alloc} + size_{alloc} - 1]$: the pointer used is not a reaching definition of the pointer allocated at n_{alloc} (negation of property (iii)).

Oracle

```

1  int *p=malloc(4);
2  p_alias=p;
3  free(p);
4  if(cond){
5    p=malloc()
6  }
7  else{
8    p=malloc();
9    p_alias=p;
10 }
11 *p=0; // never uaf
12 *p_alias=0; // uaf if cond

```

First path

- Path predicate:

$$\varphi = (p_0 = S_{alloc} \wedge p_alias_0 = p_0 \wedge p_1 = 0 \times 8040000)$$
- $\Phi' = (p_0 \neq S_{alloc} \vee p_alias_0 \notin [S_{alloc}, S_{alloc} + 3])$
- $\varphi \wedge \Phi'$ UNSAT: there is a Use-After-Free.

Oracle

```
1  int *p=malloc(4);
2  p_alias=p;
3  free(p);
4  if(cond){
5    p=malloc()
6  }
7  else{
8    p=malloc();
9    p_alias=p;
10 }
11 *p=0; // never uaf
12 *p_alias=0; // uaf if cond
```

Second path

- Path predicate: $\varphi = (p_0 = S_{alloc} \wedge p_alias_0 = p_0 \wedge p_1 = 0x8040000) \wedge p_alias_1 = p_1)$
- $\Phi' = (p_1 \neq S_{alloc} \vee p_alias_1 \notin [S_{alloc}, S_{alloc} + 3])$
- $\varphi \wedge \Phi'$ SAT (e.g., $S_{alloc} = 0$), no Use-After-Free.

Initial Memory

```
1 char key[]="ABC";
2
3 int f(int i){ // 0 <= i < 3
4     if(key[i] == 'B'){
5         // dest to reach
6     }
7 }
```

$$\Phi_{t_4} \triangleq 0 \leq i < 3 \wedge \text{key}[i] \neq B'.$$

The inversion of the last condition leads then to:

$$\Phi'_{t_4} \triangleq 0 \leq i < 3 \wedge \text{key}[i] = B'.$$

A solution: $i = 2$; $\text{key}[2] = B'$. Yet $\text{key}[2]$ is not user-controllable.

Initial Memory

Possible solutions

- Let a free (symbolic) uninitialized memory → lost of the correctness;
- Initial state: a snapshot → heavy for the solver;
- User define the initial state → complex use;
- If all read are concretized → concretizing every byte read before being written

Initial Memory

Our solution

- (i) We consider as *valid* only models with constraints on symbolic variables corresponding to inputs;
- (ii) We refine the path predicate if it generates an invalid model;
- Refine = re-executing the trace and gathering concrete values;
- Recursively until a valid model is found (or *UNSAT*);
- Necessary for JasPer,
- Possible improvements (gathering several values in one round, ...).

Libraries

Handling libraries

- Libraries are widely present in real-world;
- Do not want to explore all of them;
- We use two solutions:
 - Stubs → model effects on path predicate without tracing instructions (e.g.: if *realloc* returns new pointer, it performs copy of data);
 - Library Driven Heuristics (LDH) → known-behavior used to improve guiding.

Library Driven Heuristics (LDH): Example

```
1  p=malloc(size) ;
2  if(p == NULL)
3  {
4      // path to trigger
5  }
```

Example

- To trigger the path → *malloc* needs to return 0.

malloc heuristic

- Principle: on allocation functions, try as parameter a large number.

Library Driven Heuristics (LDH): Example

```
1  read(f,tmp,255);
2  for(i=0;i<255;i++){
3      if(tmp[i]=='\0') break;
4      buf[i] = tmp[i];
5  }
6  buf[i]='\0';
7  if(strcmp(buf,"this is really bad") == 0)
8      ..
```

Example

- Comparison on string whose length depends of loop iteration number;
- Here, every time loop iterates: $buf[i] \neq '\0'$;
- To solve *strcmp* condition, need a trace unrolling 18 times the loop.

Library Driven Heuristics (LDH): Example

```
1  read(f,tmp,255);
2  for(i=0;i<255;i++){
3      if(tmp[i]=='\0') break;
4      buf[i] = tmp[i];
5  }
6  buf[i]='\0';
7  if(strcmp(buf,"this is really bad") == 0)
8      ..
```

strcmp heuristic

- Principle: on this pattern, use size of constant strings passed to *strcmp* to find the desired iteration.


CVE on JasPer

```
MIF
component
```

Figure: PoC of CVE-2015-5221 (test.plain)

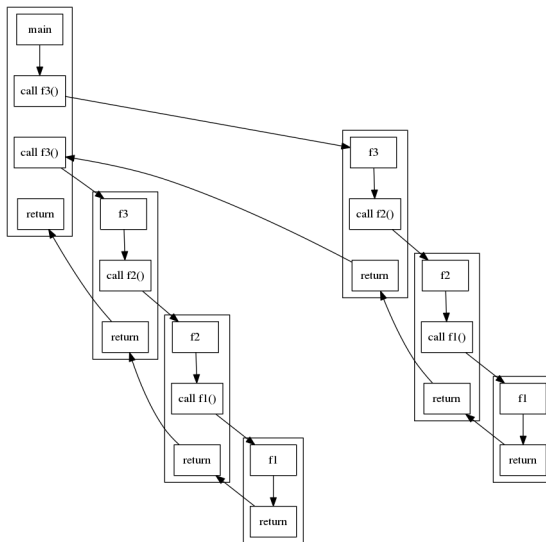
```
jasper --input test.plain --input-format mif --output out --output-format mif
```

Static Analysis Results

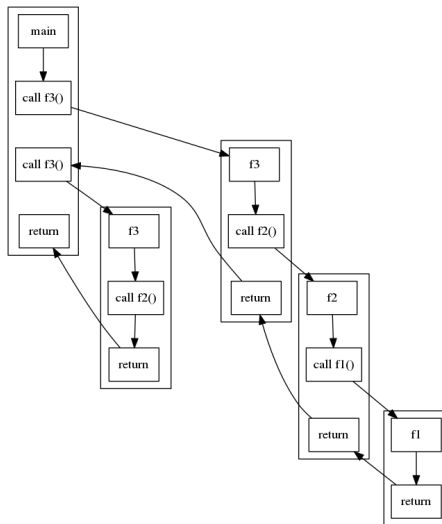


Inlining

```
1  void f1(){
2    return ;
3  }
4
5  void f2(){
6    f1();
7    return ;
8  }
9
10 void f3(){
11    f2();
12    return ;
13 }
14
15 void main(){
16    f3();
17    f3();
18 }
```



Inlining bounded by size (4)



Inlining bounded by depth (2)

